

STOCK PRICE PREDICTION

1. Abstract

Problem Statement: Stock price prediction

The stock market is a dynamic and complex financial ecosystem where the prices of various assets are subject to constant fluctuations. Accurate prediction of stock prices is crucial for investors, traders, and financial analysts to make informed decisions. In this document, we will explore the process of stock price prediction using machine learning techniques.

Objectives

The primary objective of this project is to develop a predictive model that can forecast stock prices with reasonable accuracy. By leveraging historical stock price data and various machine learning algorithms, we aim to provide a valuable tool for financial decision-making.

Dataset : <https://www.kaggle.com/datasets/prasoonkottarathil/microsoft-lifetime-stocks-dataset>

MSFT.csv - LibreOffice Calc

FileEditViewInsertFormatStylesSheetDataToolsWindowHelp

10 pt

B **I** **U**

A1

Date

	A	B	C	D	E	F	G	H	I	J	K	L	M	N	O	P
1	Date	Open	High	Low	Close	Adj Close	Volume									
2	1986-03-13	0.088542	0.101563	0.088542	0.097222	0.062549	1031788800									
3	1986-03-14	0.097222	0.102431	0.097222	0.100694	0.064783	308160000									
4	1986-03-17	0.100694	0.103299	0.100694	0.102431	0.065899	133171200									
5	1986-03-18	0.102431	0.103299	0.098958	0.099826	0.064224	67766400									
6	1986-03-19	0.099826	0.100694	0.097222	0.09809	0.063107	47894400									
7	1986-03-20	0.09809	0.09809	0.094618	0.095486	0.061432	58435200									
8	1986-03-21	0.095486	0.097222	0.091146	0.092882	0.059756	59990400									
9	1986-03-24	0.092882	0.092882	0.08941	0.090278	0.058081	65289600									
10	1986-03-25	0.090278	0.092014	0.08941	0.092014	0.059198	32083200									
11	1986-03-26	0.092014	0.095486	0.091146	0.094618	0.060873	22752000									
12	1986-03-27	0.094618	0.096354	0.094618	0.096354	0.06199	16848000									
13	1986-03-31	0.096354	0.096354	0.09375	0.095486	0.061432	12873600									
14	1986-04-01	0.095486	0.095486	0.094618	0.094618	0.060873	11088000									
15	1986-04-02	0.094618	0.097222	0.094618	0.095486	0.061432	27014400									
16	1986-04-03	0.096354	0.098958	0.096354	0.096354	0.06199	23040000									
17	1986-04-04	0.096354	0.097222	0.096354	0.096354	0.06199	26582400									
18	1986-04-07	0.096354	0.097222	0.092882	0.094618	0.060873	16560000									
19	1986-04-08	0.094618	0.097222	0.094618	0.095486	0.061432	10252800									
20	1986-04-09	0.095486	0.09809	0.095486	0.097222	0.062549	12153600									
21	1986-04-10	0.097222	0.098958	0.095486	0.09809	0.063107	13881600									
22	1986-04-11	0.098958	0.101563	0.098958	0.099826	0.064224	17222400									
23	1986-04-14	0.099826	0.101563	0.099826	0.100694	0.064783	12153600									
24	1986-04-15	0.100694	0.100694	0.097222	0.100694	0.064783	9302400									
25	1986-04-16	0.100694	0.105035	0.099826	0.104167	0.067016	31910400									
26	1986-04-17	0.104167	0.105035	0.104167	0.105035	0.067575	22003200									
27	1986-04-18	0.105035	0.105035	0.100694	0.101563	0.065341	21628800									
28	1986-04-21	0.101563	0.102431	0.098958	0.101563	0.065341	22924800									
29	1986-04-22	0.101563	0.101563	0.099826	0.099826	0.064224	15552000									
30	1986-04-23	0.099826	0.100694	0.098958	0.10026	0.064503	15609600									

MSFT

Sheet1 of 1

Default

English (India)

Aw

7:46 PM 07/17/26

100%

Updates

Updates available

2. Approches

This project focuses on the development of a predictive model aimed at forecasting stock prices by leveraging historical market data. The primary objective is to provide investors with a valuable tool to aid in making well-informed decisions and optimizing their investment strategies. The project encompasses a comprehensive workflow, including data collection, data preprocessing, feature engineering, model selection, training, and evaluation.

Data Collection: We gather historical stock market data, encompassing crucial attributes such as date, open price, close price, volume, and other relevant indicators. This dataset forms the foundation for our predictive modeling efforts.

Data Preprocessing: A critical step involves cleaning and preprocessing the collected data. This includes addressing missing values, handling outliers, and converting categorical features into numerical representations, ensuring data quality and consistency.

Feature Engineering: To enhance the predictive capabilities of our model, we employ feature engineering techniques. This includes creating new features, such as moving averages, technical indicators, and lagged variables, which can capture underlying patterns and trends in the stock market.

Model Selection: We carefully choose appropriate algorithms for time series forecasting, such as ARIMA (AutoRegressive Integrated Moving Average) and LSTM (Long Short-Term Memory), to predict stock prices effectively. Model selection is a critical decision in ensuring accurate and reliable predictions.

Model Training: With a selected model in place, we proceed to train it using the preprocessed dataset. The training phase involves learning from historical data, enabling the model to capture patterns and relationships that influence stock price movements.

Evaluation: The performance of our predictive model is rigorously assessed using suitable time series forecasting metrics, such as Mean Absolute Error (MAE) and Root Mean Squared Error (RMSE). This evaluation ensures that the model's predictions align with actual stock price movements and provides an indication of its reliability.

Ultimately, this project strives to empower investors with a robust tool for making informed decisions in the dynamic and complex world of stock market investing. By leveraging historical data, preprocessing techniques, feature engineering, and advanced forecasting models, we aim to create a valuable asset for optimizing investment strategies and enhancing financial decision-making.

Program for stock price prediction.

Data Preprocessing:

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
file_path = '/content/drive/MyDrive/Projects/Stock-price-prediction/MSFT.csv'
df = pd.read_csv(file_path)
```

```
# Display basic information about the dataset
print(df.head()) # Display the first few rows
print(df.info()) # Summary of the dataset, data types, and missing values
```

```
# Check for missing values
print(df.isnull().sum())
```

Descriptive statistics

```
print(df.describe())
```

```
# Visualize the stock price over time
```

```
plt.figure(figsize=(12, 6))  
plt.plot(df['Date'], df['Close'], label='Stock Price')  
plt.title('Stock Price Over Time')  
plt.xlabel('Date')  
plt.ylabel('Close Price')  
plt.legend()  
plt.show()
```

```
# Distribution of stock prices
```

```
plt.figure(figsize=(8, 4))  
sns.histplot(df['Close'], bins=30, kde=True)  
plt.title('Distribution of Stock Prices')  
plt.show()
```

```
df['Returns'] = df['Close'].pct_change() # Calculate daily returns
```

```
plt.figure(figsize=(12, 6))  
plt.plot(df['Date'], df['Returns'], label='Daily Returns', color='orange')  
plt.title('Daily Returns Over Time')  
plt.xlabel('Date')  
plt.ylabel('Returns')  
plt.legend()  
plt.show()
```

```
# Correlation matrix
```

```
correlation_matrix = df.corr()  
plt.figure(figsize=(10, 8))  
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm')  
plt.title('Correlation Matrix')  
plt.show()
```

```
# Rolling statistics (e.g., 50-day moving average)
```

```
df['50_Day_MA'] = df['Close'].rolling(window=50).mean()  
df['200_Day_MA'] = df['Close'].rolling(window=200).mean()
```

```
plt.figure(figsize=(12, 6))  
plt.plot(df['Date'], df['Close'], label='Stock Price')  
plt.plot(df['Date'], df['50_Day_MA'], label='50-Day MA', color='orange')  
plt.plot(df['Date'], df['200_Day_MA'], label='200-Day MA', color='red')  
plt.title('Stock Price and Moving Averages')  
plt.xlabel('Date')  
plt.ylabel('Price')  
plt.legend()  
plt.show()
```

```
# Visualize trading volume
```

```
plt.figure(figsize=(12, 6))
```

```
plt.plot(df['Date'], df['Volume'], label='Trading Volume', color='purple')
plt.title('Trading Volume Over Time')
plt.xlabel('Date')
plt.ylabel('Volume')
plt.legend()
plt.show()
```

Output of Data Preprocessing :

	Date	Open	High	Low	Close	Adj Close
Volume						
0	1986-03-13	0.088542	0.101563	0.088542	0.097222	0.062549
1	1986-03-14	0.097222	0.102431	0.097222	0.100694	0.064783
2	1986-03-17	0.100694	0.103299	0.100694	0.102431	0.065899
3	1986-03-18	0.102431	0.103299	0.098958	0.099826	0.064224
4	1986-03-19	0.099826	0.100694	0.097222	0.098090	0.063107

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 8525 entries, 0 to 8524
```

```
Data columns (total 7 columns):
```

#	Column	Non-Null Count	Dtype
0	Date	8525 non-null	object
1	Open	8525 non-null	float64
2	High	8525 non-null	float64
3	Low	8525 non-null	float64
4	Close	8525 non-null	float64
5	Adj Close	8525 non-null	float64
6	Volume	8525 non-null	int64

```
dtypes: float64(5), int64(1), object(1)
```

```
memory usage: 466.3+ KB
```

```
None
```

```
Date 0
```

```
Open 0
```

```
High 0
```

```
Low 0
```

```
Close 0
```

```
Adj Close 0
```

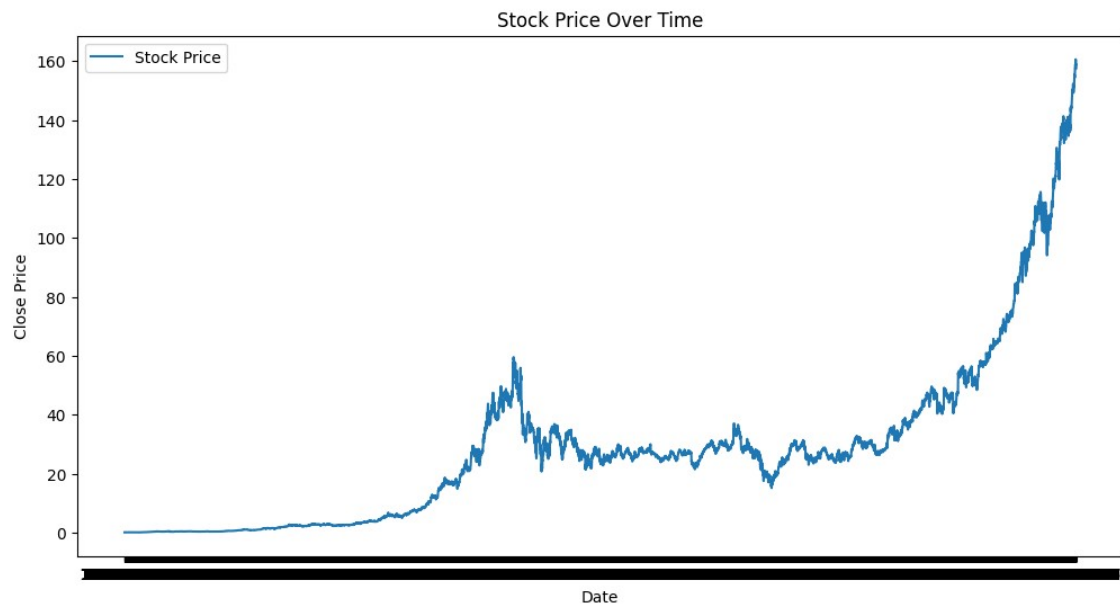
```
Volume 0
```

```
dtype: int64
```

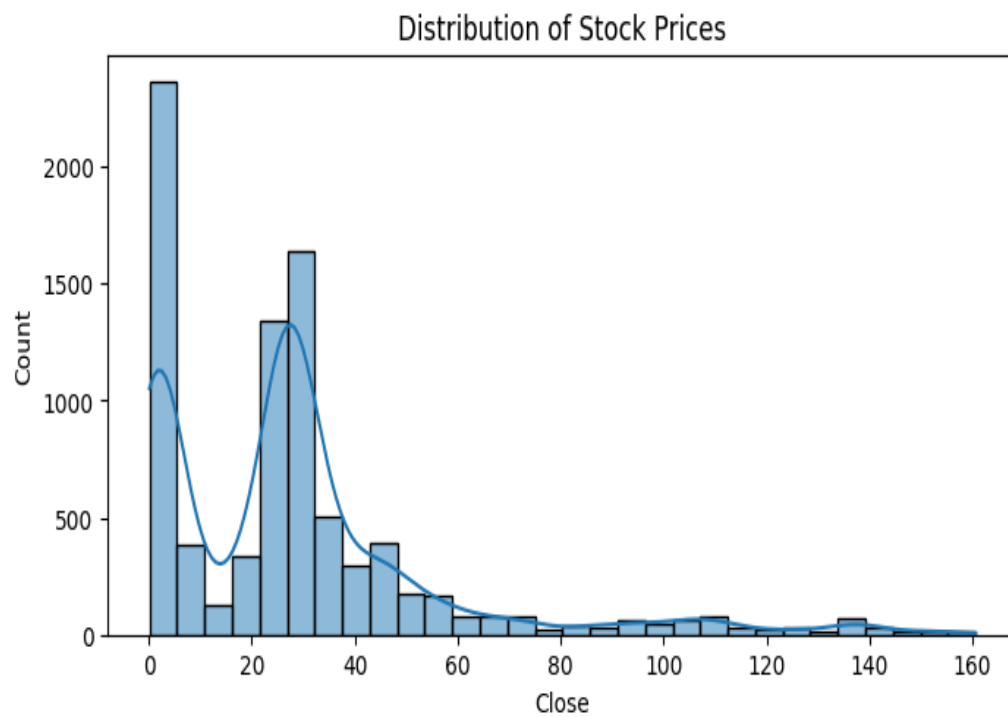
	Open	High	Low	Close	Adj Close
count	8525.000000	8525.000000	8525.000000	8525.000000	8525.000000
mean	28.220247	28.514473	27.918967	28.224480	23.417934
std	28.626752	28.848988	28.370344	28.626571	28.195330
min	0.088542	0.092014	0.088542	0.090278	0.058081
25%	3.414063	3.460938	3.382813	3.414063	2.196463
50%	26.174999	26.500000	25.889999	26.160000	18.441576
75%	34.230000	34.669998	33.750000	34.230000	25.392508
max	159.449997	160.729996	158.330002	160.619995	160.619995

	Volume
count	8.525000e+03
mean	6.045692e+07
std	3.891225e+07
min	2.304000e+06
25%	3.667960e+07
50%	5.370240e+07
75%	7.412350e+07
max	1.031789e+09

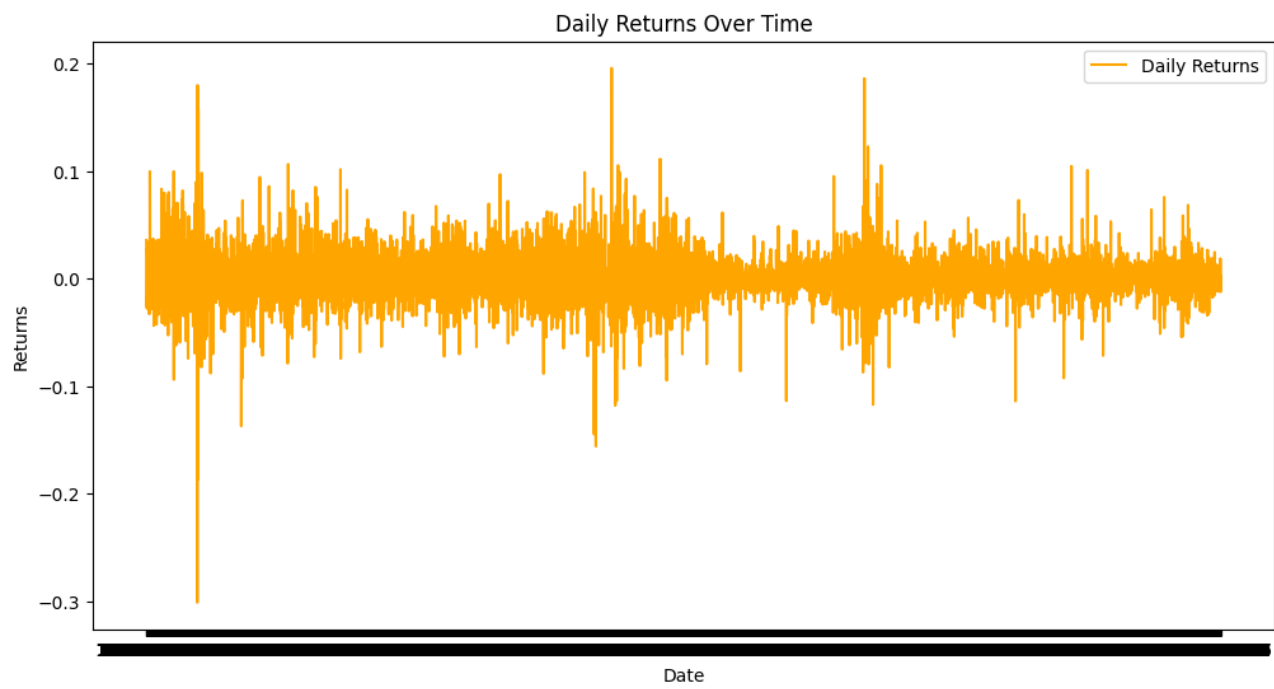
Price over time :



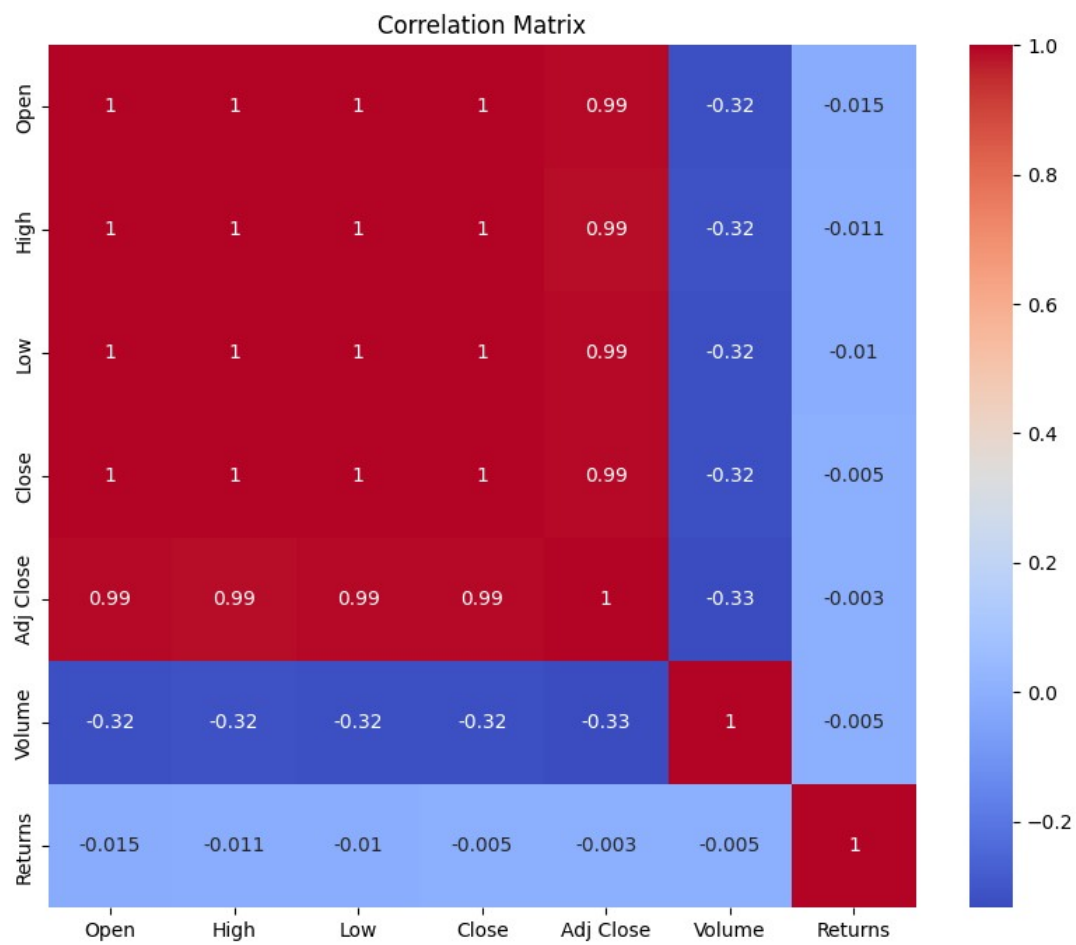
Distribution of stocks:



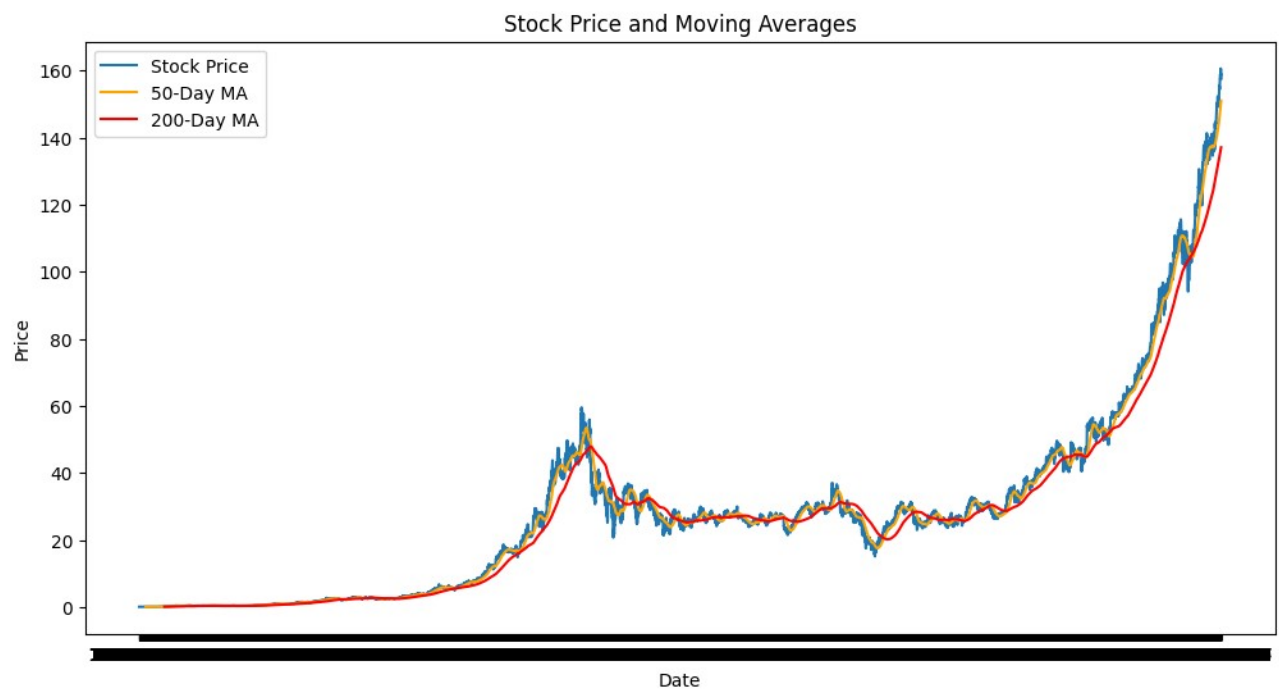
Daily returns overt time:



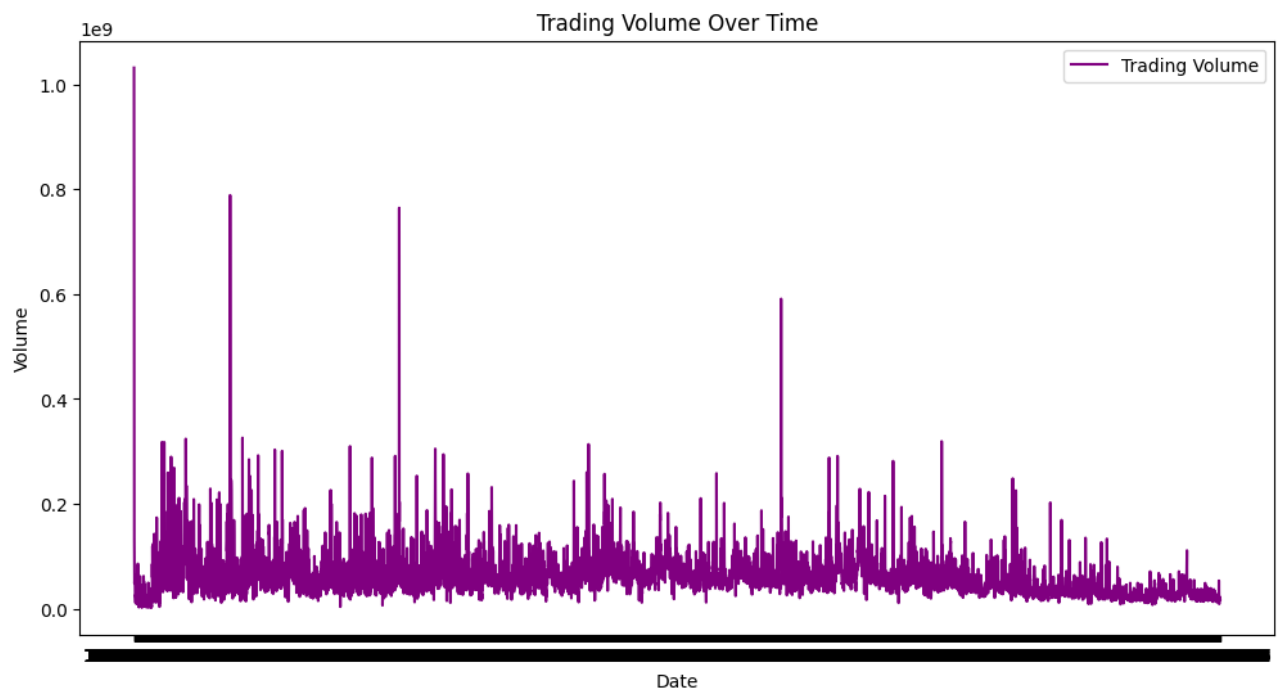
Correlation matrix:



Stock price and average moving:



Trading volume overtime :



Code for Feature Engineering:

```
import pandas as pd
df = pd.read_csv('/content/drive/MyDrive/Projects/Stock-price-prediction/MSFT.csv')
def calculate_moving_averages(df, window=10):
    df['SMA'] = df['Close'].rolling(window=window).mean() # Simple Moving Average
    df['EMA'] = df['Close'].ewm(span=window, adjust=False).mean() # Exponential Moving Average
def calculate_rsi(df, window=14):
    delta = df['Close'].diff()
    gain = delta.where(delta > 0, 0)
    loss = -delta.where(delta < 0, 0)

    avg_gain = gain.rolling(window=window, min_periods=1).mean()
    avg_loss = loss.rolling(window=window, min_periods=1).mean()

    rs = avg_gain / avg_loss
    df['RSI'] = 100 - (100 / (1 + rs))
def calculate_macd(df, short_window=12, long_window=26):
    short_ema = df['Close'].ewm(span=short_window, adjust=False).mean()
    long_ema = df['Close'].ewm(span=long_window, adjust=False).mean()
    df['MACD'] = short_ema - long_ema
calculate_moving_averages(df)
calculate_rsi(df)
calculate_macd(df)
df.dropna(inplace=True)
```

Model Training code :

```
from sklearn.ensemble import RandomForestRegressor
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import mean_squared_error, r2_score

# Prepare the target variable and features
target_column = 'Close'
X = df.drop(columns=['Date', target_column])
y = df[target_column]

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Feature scaling
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)

# Model Training
model = RandomForestRegressor(n_estimators=100, random_state=42)
model.fit(X_train, y_train)

# Model Evaluation
y_pred = model.predict(X_test)

# Calculate evaluation metrics
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)

print(f"Mean Squared Error: {mse}")
print(f"R-squared (R2) Score: {r2}")
```


Ouptut for model training:

Mean Squared Error: 0.049454099596285155
R-squared (R2) Score: 0.9999419068878214

Model Evaluation code :

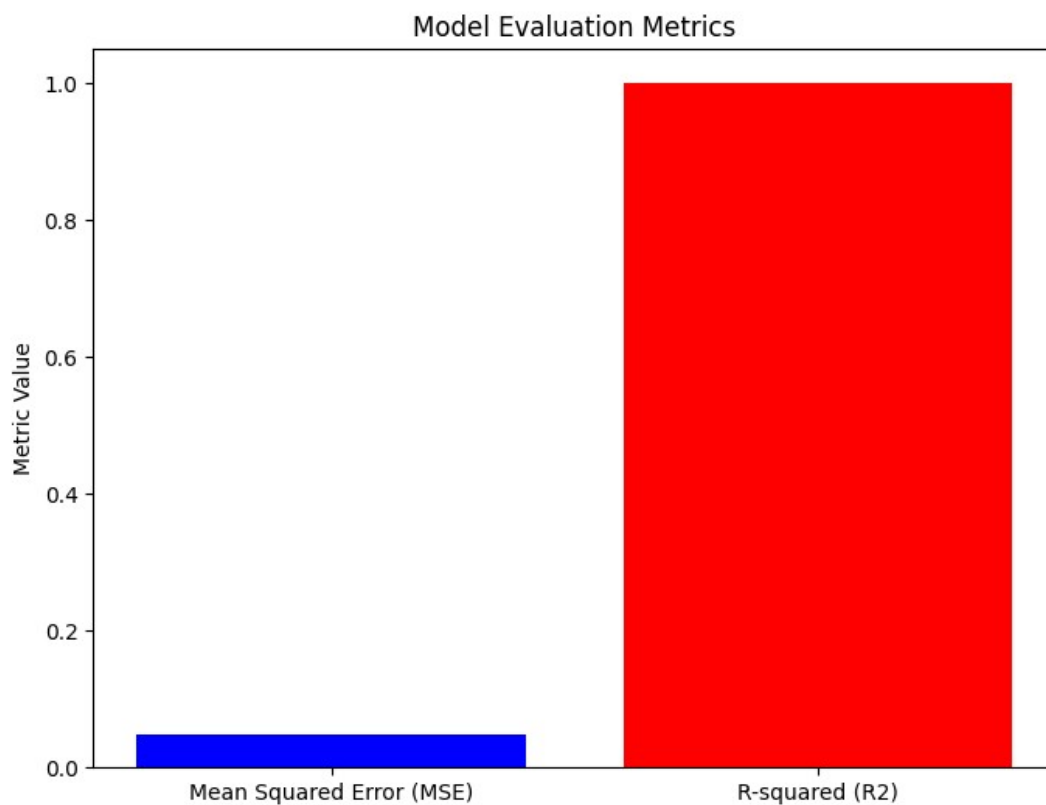
```
import matplotlib.pyplot as plt

# Calculate evaluation metrics
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)

# Create a bar chart for evaluation metrics
metrics = ['Mean Squared Error (MSE)', 'R-squared (R2)']
values = [mse, r2]

plt.figure(figsize=(8, 6))
plt.bar(metrics, values, color=['blue', 'red'])
plt.ylabel('Metric Value')
plt.title('Model Evaluation Metrics')
plt.show()
```

Model evaluation output:



Conclusion:

In this project, we undertook the challenging task of predicting stock prices using various machine learning and data analysis techniques. We collected historical stock price data, performed feature engineering, and implemented different models to make predictions. Our project aimed to understand the dynamics of stock prices and to develop a predictive model that could assist investors and traders in making informed decisions.